



Coding.net 轻量级运维实践

吴柯

What ?
Coding.net 是什么 ?



关于Coding.net

Coding.net 是一个面向开发者的云端开发协作平台



项目协作



代码托管



运行空间



质量控制



自动测试




持续集成



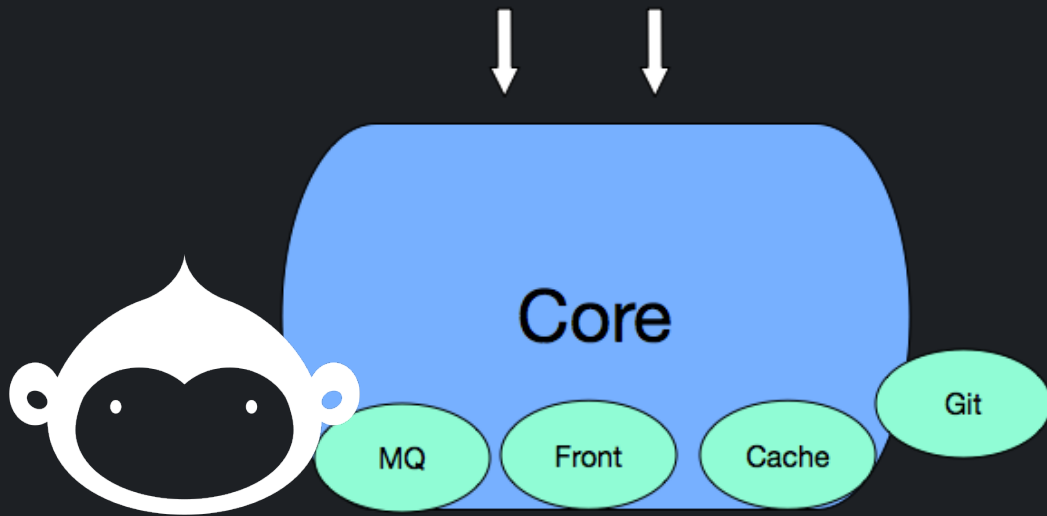
在线IDE

Coding v0.1时代

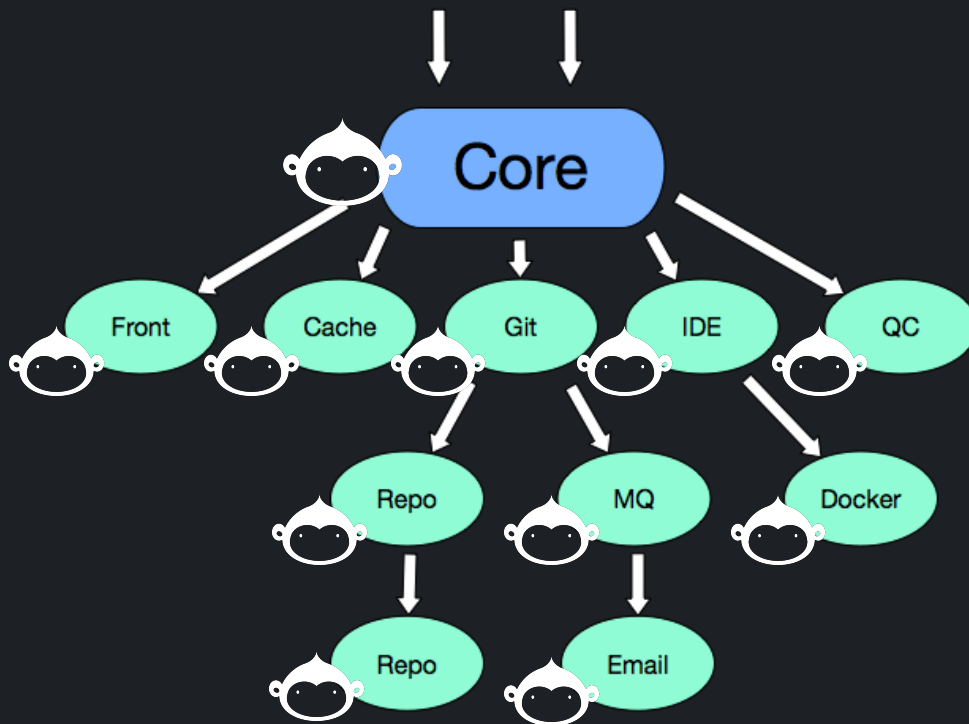
- 保证开发速度，服务模块过于耦合
- 部署上线很虐心
- 运维变更靠超人
- 监控靠有限的外部服务
- 救火常态化



一开始..



后来..



Coding v1.0

- 服务模块化
- 通信接口化
- 部署脚本化
- 主机管理批量化



Restful API

- GET
- POST
- PUT
- DELETE

账户			账户
/account/check	GET		- 动态
/account/register	POST		社交
/account/login	POST		- 冒泡
/account/logout	POST		项目
/account/captcha/{action}	GET		- 列表
/account/activate	GET	POST	- 讨论
/account/reset_password	GET	POST	- 文档
/account/invite	POST	GET	- 任务
/account/invite/register	POST		代码托管 Git
/account/name/{username}	GET		- Pull Request
/account/name/{username}/no_cache	GET		- Merge Request
/account/current_user	GET		演示部署平台
/account/avatar	GET		Markdown解析
/account/search	GET		
/account/update_info	POST		
/account/get_notice_settings	GET		
/account/change_notice_setting	GET		
/account/update_pwd	POST		
/account/projects/pin	GET	POST	DELETE



- 基于SSH管理
- 批量管理&配置主机
- inventory 主机分组管理
- 支持Playbook

/hosts file:

[backup]

backup-1.coding.local ansible_ssh_user=root

[core-app]

core-app-1.coding.local ansible_ssh_user=root ansible_ssh_port=2222

core-app-2.coding.local ansible_ssh_user=root ansible_ssh_port=2222

[core-db]

core-db-1.coding.local ansible_ssh_user=root

core-db-2.coding.local ansible_ssh_user=root

[core-cache]

core-cache-1.coding.local ansible_ssh_user=root

core-cache-2.coding.local ansible_ssh_user=root

[core-queue]

core-queue-1.coding.local ansible_ssh_user=root

core-queue-2.coding.local ansible_ssh_user=root

[core:children]

core-app

core-db

core-cache

core-queue

批量远程操作

```
~$ ansible 'core' -i /home/ubuntu/ansible/hosts -a 'uptime'
```

```
core-cache.coding.local | success | rc=0 >>
```

```
17:50:56 up 152 days, 38 min, 1 user, load average: 0.30, 0.37, 0.36
```

```
core-db.coding.local | success | rc=0 >>
```

```
17:50:56 up 152 days, 33 min, 1 user, load average: 0.23, 0.23, 0.18
```

```
core-queue.coding.local | success | rc=0 >>
```

```
17:50:56 up 152 days, 37 min, 1 user, load average: 0.03, 0.04, 0.05
```

```
core-app.coding.local | success | rc=0 >>
```

```
17:50:56 up 152 days, 36 min, 1 user, load average: 0.99, 0.89, 0.95
```

```
.....
```

```
~$
```

Puppet & Ansible



Puppet

- 支持主流的操作系统
- 历史悠久相对成熟
- 基于Ruby 性能相对慢

Ansible

- 基于SSH 无代理
- 安装简单
- 命令行操作简单方便
- 基于python 速度快，性能好


但是...

- 60%的批量操作是配置环境
- 众多微服务难管理
- 自动化脚本难维护
- 监控体系不健全
- 日志分散数据沉睡




Coding v2.0

- 服务容器化
- 操作代码化
- 部署自动化
- 建立监控体系
- 日志收集&分析




容器化过程



运行代码化

```
jobs : <
  name: "coding-front"
  image: "coding-front:X"
>
jobs : <
  name: "coding-backend"
  image: "coding-backend:Y"
  env: <
    key: "xxxxxxx"
    value: "xxxxxxx"
  >
>
jobs: <
  name: "mysql"
  image: "mysql:5.5"
>
```




运行代码化

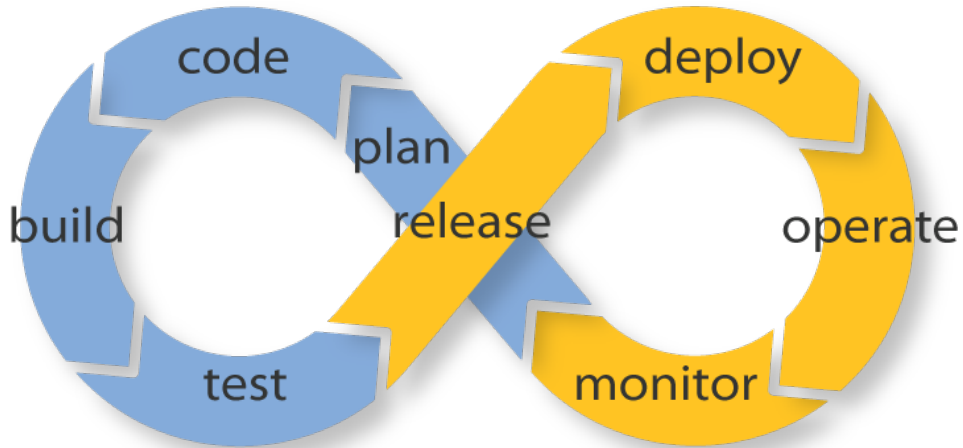
```
$ go run stack.go up  
Job: coding-front  
Image: coding-front:X  
State: [/coding-front_X]: Up 1s
```

```
Job: coding-backend  
Image: coding-backend:Y  
State: [/coding-backend_Y]: Up 1s
```

```
Job: mysql  
Image: mysql:5.5  
State: [/mysql_5.5]: Up 1s
```

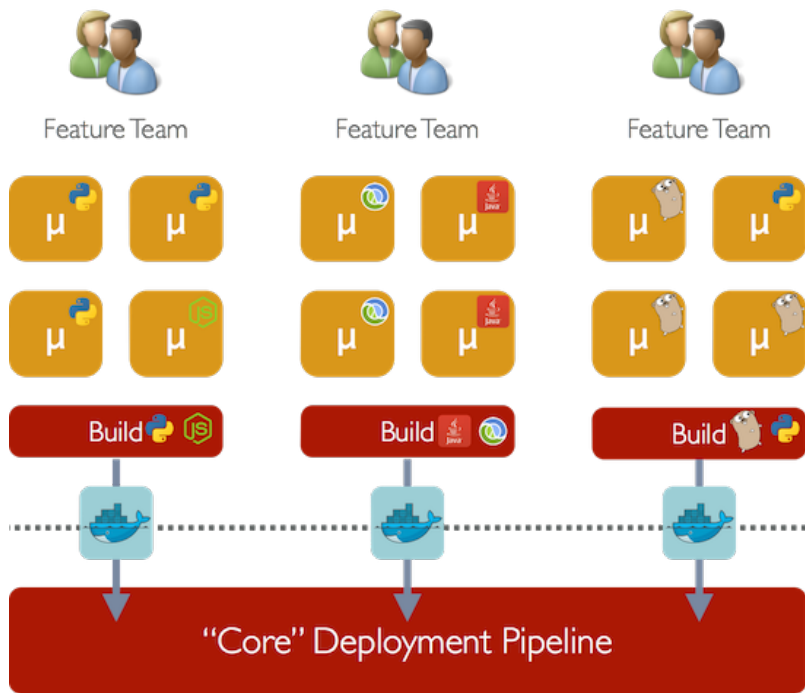


DevOps

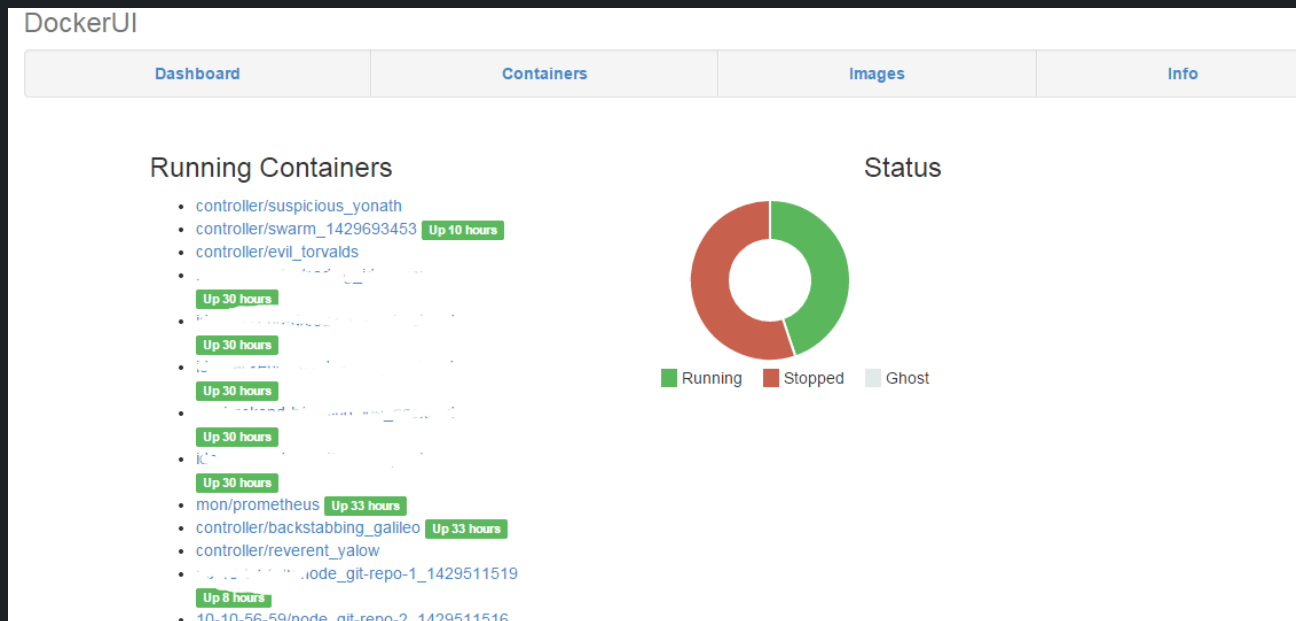


Endless Possibilities: DevOps can create an infinite loop of release and feedback for all your code and deployment targets.

部署：Pipeline 模式



部署 Docker UI



部署 Docker UI

DockerUI

[Dashboard](#)[Containers](#)[Images](#)[Info](#)

Container: /swarm_1429693453

[Rename](#)[Stop](#)[Kill](#)[Pause](#)

Created: 2015-04-22T09:04:13.66923639Z

Path: swarm

Args: ["manage", "file:///etc/swarm/hosts.swarm"]

Exposed Ports:

- 2375/tcp
- 9999/tcp

Environment:

- SWARM_HOST=:9999
- PATH=/go/bin:/usr/src/go/bin:/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
- GOLANG_VERSION=1.3.3
- GOPATH=/go/src/github.com/docker/swarm/Godeps/_workspace:/go
- DEBIAN_FRONTEND=noninteractive
- LANG=en_US.UTF-8
- LANGUAGE=en_US:en
- LC_ALL=en_US.UTF-8

Publish All: false

Ports:

Hostname: controller

部署 Docker UI

DockerUI

Dashboard

Containers

Images

Info

Logs for container: /swarm_1429693453

stdout

slderr

Reload logs

200

lines

☐ Timestamps

STDOUT

STDERR

```
time="2015-04-22T18:47:58Z" level=error msg="Get http://paas-service-mongodb-2.coding.local:2375/v1.15/info: dial tcp 10.9.27.25:2375: connection refused"
time="2015-04-22T18:47:58Z" level=error msg="Get http://paas-service-postgresql-1.coding.local:2375/v1.15/info: dial tcp 10.9.27.243:2375: connection refused"
time="2015-04-22T18:47:58Z" level=error msg="Get http://paas-service-redis-1.coding.local:2375/v1.15/info: dial tcp 10.9.21.212:2375: connection refused"
time="2015-04-22T18:47:58Z" level=error msg="Get http://paas-service-db-1.coding.local:2375/v1.15/info: dial tcp 10.9.1.188:2375: connection refused"

time="2015-04-22T18:47:58Z" level=error msg="Get http://paas-container-3.coding.local:2375/v1.15/info: dial tcp 10.9.19.86:2375: connection refused"
time="2015-04-22T18:47:58Z" level=error msg="Get http://paas-container-6.coding.local:2375/v1.15/info: dial tcp 10.9.24.110:2375: connection refused"

time="2015-04-22T18:47:58Z" level=error msg="Get http://paas-container-2.coding.local:2375/v1.15/info: dial tcp 10.9.19.61:2375: connection refused"
time="2015-04-22T18:47:58Z" level=error msg="Get http://paas-container-12.coding.local:2375/v1.15/info: dial tcp 10.9.31.236:2375: connection refused"
"
time="2015-04-22T18:47:58Z" level=error msg="Get http://qc-db-1.coding.local:2375/v1.15/info: dial tcp 10.10.31.77:2375: connection refused"
```

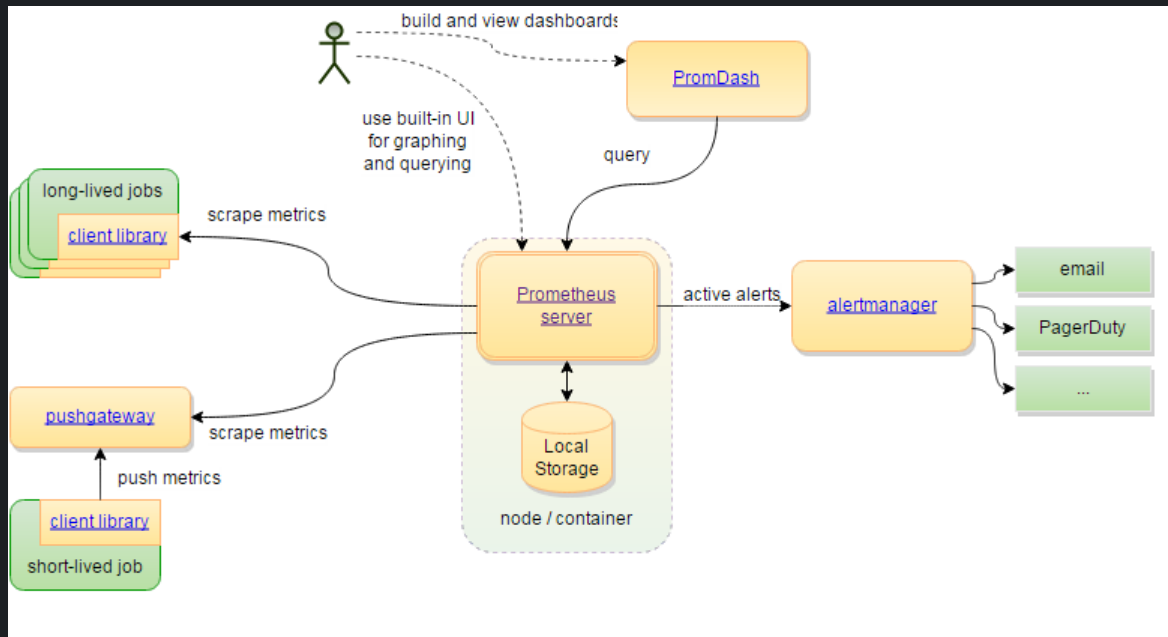
监控体系

Prometheus 是SoundCloud开发的一个开源的系统
监控和报警工具

- 多维度数据模型存储 (基于时间序列的 key/value)
- 灵活简单的多维度查询语句
- 不依赖于复杂的分布式的存储模型，单节点服务即装即用
- 数据收集可以通过 HTTP 拉取，也可以通过 Gateway 推送
- 静态配置和动态发现监控目标
- 支持多种主流的图形和仪表模型




监控体系



Prometheus & InfluxDB



- 存储事件的所有元数据
- 数据占用空间大
- 分布式集群设计，部署管理复杂
- 适用于大型数据处理




- 只存储数字化的数据
- 数据占用空间小
- 单节点即可运行，部署管理简单


监控体系

Node Overview - core-app-1.coding.local:9100

CPU Usage



Disk I/O Utilization



Overview

User CPU	2.4%
System CPU	0.8%
Memory Total	15.67GiB
Memory Free	3.524GiB

Network

docker0 Received	0B/s
docker0 Transmitted	0B/s
eth0 Received	159.5kB/s
eth0 Transmitted	157kB/s

Disks

vda Utilization	0.1%
vdb Utilization	0.0%
vda Throughput	27.12kB/s
vdb Throughput	7.691kB/s

Filesystem Fullness

/	78.9 %
/etc/hostname	78.9 %

配置代码化

master coding-ops / prod / monitoring / conf / consoles / node-overview.html

test 8天前 21acbc0964

Update Menu, importing Nodes graphs.

file | 5.51 KB 编辑 原始数据 按行查看 提交历史 删除

```
1  {{ template "head" . }}
2
3  {{ template "prom_right_table_head" }}
4  <tr><th colspan="2">Overview</th></tr>
5  <tr>
6    <td>User CPU</td>
7    <td>{{ template "prom_query_drilldown" (args (printf "sum(rate(node_cpu{job='node',instance='%s',mode='user'})[5m])) * 100
8  </tr>
9  <tr>
10   <td>System CPU</td>
11   <td>{{ template "prom_query_drilldown" (args (printf "sum(rate(node_cpu{job='node',instance='%s',mode='system'})[5m])) * 100
12 </tr>
13 <tr>
14   <td>Memory Total</td>
15   <td>{{ template "prom_query_drilldown" (args (printf "node_memory_MemTotal{job='node',instance='%s'}" .Params.instance) "B"
16 </tr>
17 <tr>
18   <td>Memory Free</td>
19   <td>{{ template "prom_query_drilldown" (args (printf "node_memory_MemFree{job='node',instance='%s'}" .Params.instance) "B"
20 </tr>
21 <tr>
```

日志收集

Logstash conf

```
input {
  file {
    path => "/data/log/nginx/access.log"
    type => "nginx-access"
  }
}
filter {
  grok {
    type => "nginx-access"
    match => {
      "message" => "(?:(?:IPORHOST:clientip)!) - %{NUMBER:response_time:float} - (?:(?:NUMBER:response_upstream_time:float)!) - %{USER:auth} \[%{HTTPDATE:timestamp}\] \"(?:%{WORD:verb} %{NOTSPACE:request}(?: HTTP/%{NUMBER:httpversion})?|%{DATA:rawrequest})\" %{NUMBER:response:int} (?:%{NUMBER:bytes:int}!) %{QS:referrer} %{QS:agent}"
    }
  }
  date {
    type => "nginx-access"
    match => [ "timestamp" , "dd/MMM/yyyy:HH:mm:ss Z" ]
  }
}
output {
  elasticsearch {
    host => "xx.xx.xx.xx"
    protocol => transport
  }
}
```

Input Type

- file
- log4j
- eventlog
- redis
- syslog
- snmptrap
- websocket
- ...

日志收集

★ [logstash-]YYYY.MM.DD



This page lists every field in the [logstash-]YYYY.MM.DD index and the field's associated core type as recorded by Elasticsearch. While this list allows you to view the core type of each field, changing field types must be done using Elasticsearch's [Mapping API](#).

This index uses a Time-based index pattern which repeats Daily

Fields (37) Scripted Fields (0)

name	type	analyzed	indexed	popularity
_index	string	false	false	0
_type	string	false	true	1
geoip.location	geo_point	false	true	0
@version	string	false	true	0
_source	string	false	false	4
_id	string	false	false	0
request	string	true	true	11
referrer.raw	string	false	true	0
agent	string	true	true	1
auth	string	true	true	1
timestamp.raw	string	false	true	0
type	string	true	true	1
response_upstream_time	number	false	true	9
path	string	true	true	6
clientip	string	true	true	1
host	string	true	true	5
path.raw	string	false	true	0
timestamp	string	true	true	1
tags.raw	string	false	true	0
clientip.raw	string	false	true	0
message.raw	string	false	true	0
host.raw	string	false	true	0
verb	string	true	true	0
type.raw	string	false	true	0
message	string	true	true	8

分析处理

kibana Discover Visualize Dashboard Settings May 11th 2015, 14:28:09.000

Search...

[logstash-*YYYY.MM.DD]

Selected Fields

- _source

Fields

Popular fields

- bytes
- host
- message
- path
- referrer
- request
- response
- response_time
- response_upstream_time
- @timestamp
- @version
- _id
- _index
- _type
- agent
- auth
- clientip
- httpversion
- tags
- timestamp
- type
- verb

May 11th 2015, 14:28:09.000 - May 11th 2015, 14:28:10.000

Table Request Response Statistics

Elasticsearch request body

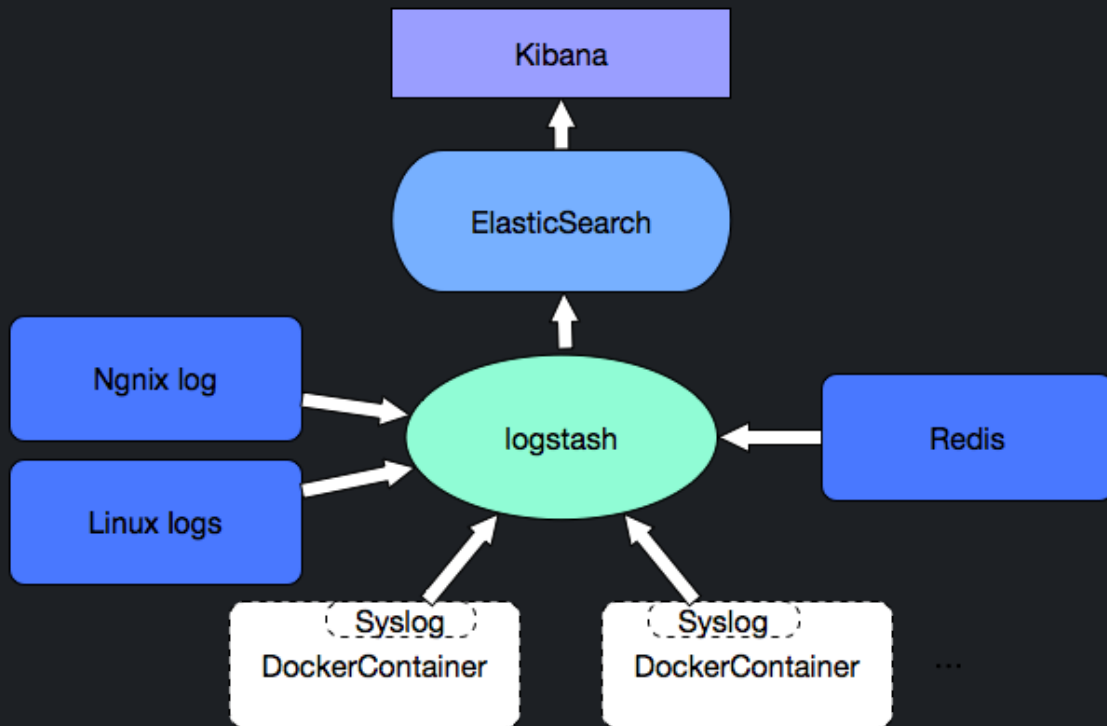
```
{
  "size": 500,
  "sort": {
    "@timestamp": "desc"
  },
  "highlight": {
    "pre_tags": [
      "@kibana-highlighted-field@"
    ],
    "post_tags": [
      "@kibana-highlighted-field@"
    ],
    "fields": {
      "*": {}
    }
  },
  "aggs": {
    "2": {
      "date_histogram": {
        "field": "@timestamp",
        "interval": "20ms",
        "pre_zone": "+08:00",
        "pre_zone_adjust_large_interval": true,
        "min_doc_count": 0,
        "extended_bounds": {
          "min": 1431325689000,
          "max": 1431325690000
        }
      }
    }
  }
}
```


Time _source

May 11th 2015, 14:28:10.000

```
message: - - 0.005 - 0.005 - - [11/May/2015:14:28:10 +0800] "POST /internal/ssh_allow HTTP/1.1" 200 213 "-" "Go 1.1 package http"
@version: 1 @timestamp: May 11th 2015, 14:28:10.000 type: nginx-access host: 10-10-58-89 path: /data/log/
response_time: 0.005 response_upstream_time: 0.005 auth: - timestamp: 11/May/2015:14:28:10 +0800 verb: POST
h_allow httpversion: 1.1 response: 200 bytes: 213 referrer: "-" agent: "Go 1.1 package http" _source: {
  "message": "5 - - [11/May/2015:14:28:10 +0800] \"POST /internal/ssh_allow HTTP/1.1\" 200 213 \\\"-\\\" \\\"Go 1.1 package http\\\""
```

日志收集&分析





wuke@coding.net